


DIABETES PREDICTION USING CLASSIFICATION

Created by:
Lin Xinyang
Shivam Varshney (456533)



OVERVIEW

1. Project description
2. Data preprocessing
3. Data visualization
4. Building models
5. performance comparison
6. further improvement.



1. Project Description



- **Objective:** To leverage Classification Algorithm to predict Diabetes in female patients.
- **Datasource:** Prima Indians Database by National Institute of Diabetes and Digestive and Kidney Diseases
- **Tools:** Five classification algorithms including Random ForestClassifier, AdaBoostClassifier, XGBClassifier, SVM, LightBGM; one regression algorithm logistic regression.

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	
5	5	116	74	0	0	25.6	0.201	30	
6	3	78	50	32	88	31.0	0.248	26	1
7	10	115	0	0	0	35.3	0.134	29	0
8	2	197	70	45	543	30.5	0.158	53	1
9	8	125	96	0	0	0.0	0.232	54	1



- **Source of the data:** The National Institute of Diabetes and Digestive and Kidney Diseases

2.Data Preprocessing

Data cleaning

```
Pregnancies      0
Glucose          0
BloodPressure    0
SkinThickness    0
Insulin          0
BMI              0
DiabetesPedigreeFunction  0
Age              0
Outcome          0
dtype: int64
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
```

```
Data columns (total 9 columns):
#   Column                                Non-Null Count
---  -----
0   Pregnancies                          768 non-null
1   Glucose                              768 non-null
2   BloodPressure                        768 non-null
3   SkinThickness                        768 non-null
4   Insulin                              768 non-null   int64
5   BMI                                  768 non-null   float64
6   DiabetesPedigreeFunction             768 non-null   float64
7   Age                                  768 non-null   int64
8   Outcome                              768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
None
```

```
# Replace missing values with column medians
```

```
diabetes_df_copy = diabetes_df.copy(deep = True)
```

```
columns_median=np.array([np.median(diabetes_df_copy["Glucose"]),
                             np.median(diabetes_df_copy["BloodPressure"]),
                             np.median(diabetes_df_copy["SkinThickness"]),
                             np.median(diabetes_df_copy["Insulin"]),
                             np.median(diabetes_df_copy["BMI"])]])
```

```
diabetes_df_copy.loc[diabetes_df_copy["Glucose"]==0,"Glucose"]= columns_median[0]
```

```
diabetes_df_copy.loc[diabetes_df_copy["BloodPressure"]==0,"BloodPressure"]= columns_median[1]
```

```
diabetes_df_copy.loc[diabetes_df_copy["SkinThickness"]==0,"SkinThickness"]= columns_median[2]
```

```
diabetes_df_copy.loc[diabetes_df_copy["Insulin"]==0,"Insulin"]= columns_median[3]
```

Data Preprocessing

```
# 1. BMI as a Range
diabetes_df_copy['BMI_Category'] = pd.cut(diabetes_df_copy['BMI'], bins=[0, 18.5, 24.9, 29.9, 34.9, 39.9],
                                          labels=['Underweight', 'Normal Weight', 'Overweight', 'Obese I', 'Obese II'])

# 2. Interaction Term: Glucose * Insulin
diabetes_df_copy['Glucose_Insulin_Interact'] = diabetes_df_copy['Glucose'] * diabetes_df_copy['Insulin']

# 3. Age Groups
diabetes_df_copy['Age_Group'] = pd.cut(diabetes_df_copy['Age'], bins=[20, 30, 40, 50, 60, 70, np.inf],
                                       labels=['20-30', '30-40', '40-50', '50-60', '60-70', '70+'])
diabetes_df_copy
```

Newly created features:




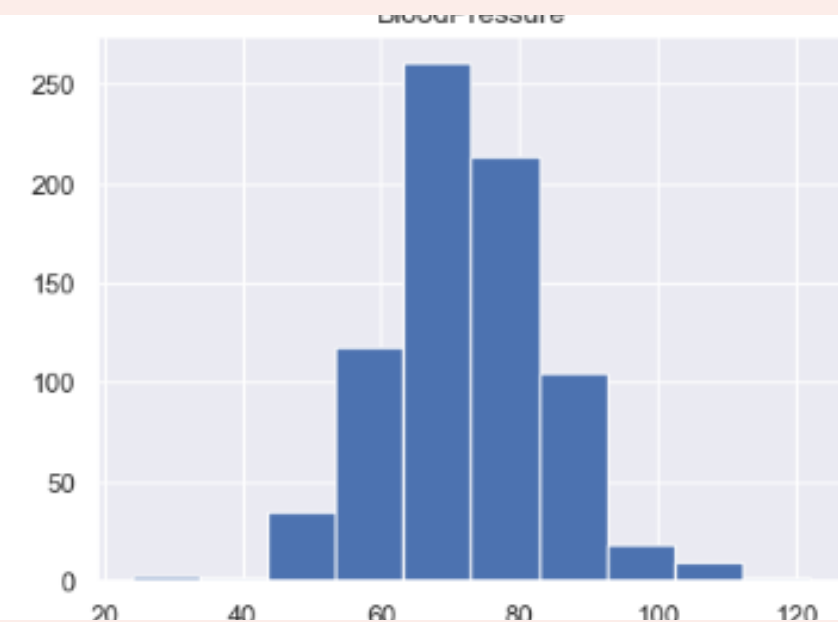

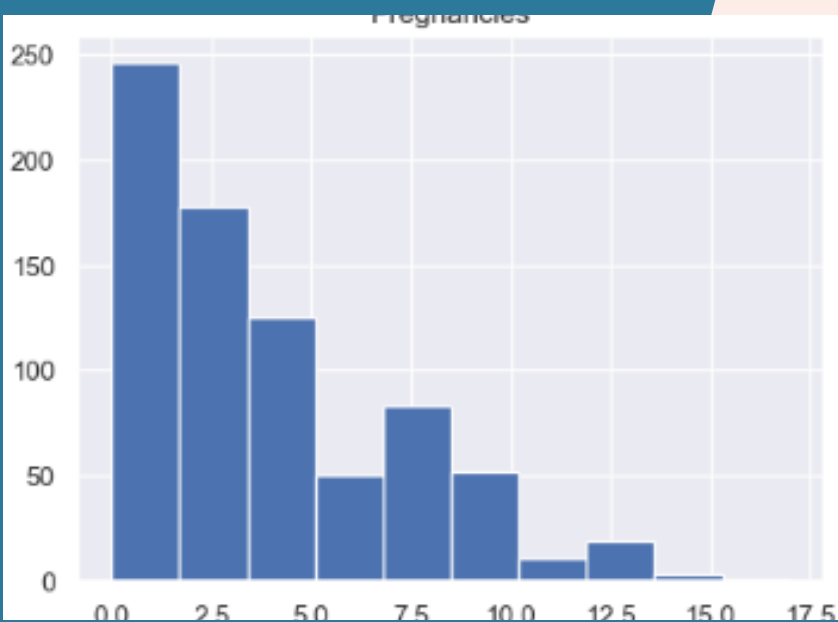



- BMI_Category (categorical)
- Age_Group (categorical)
- Glucose_Insulin_Interact (numerical)





Data
engineering

3.Data Visualization


Bar chart






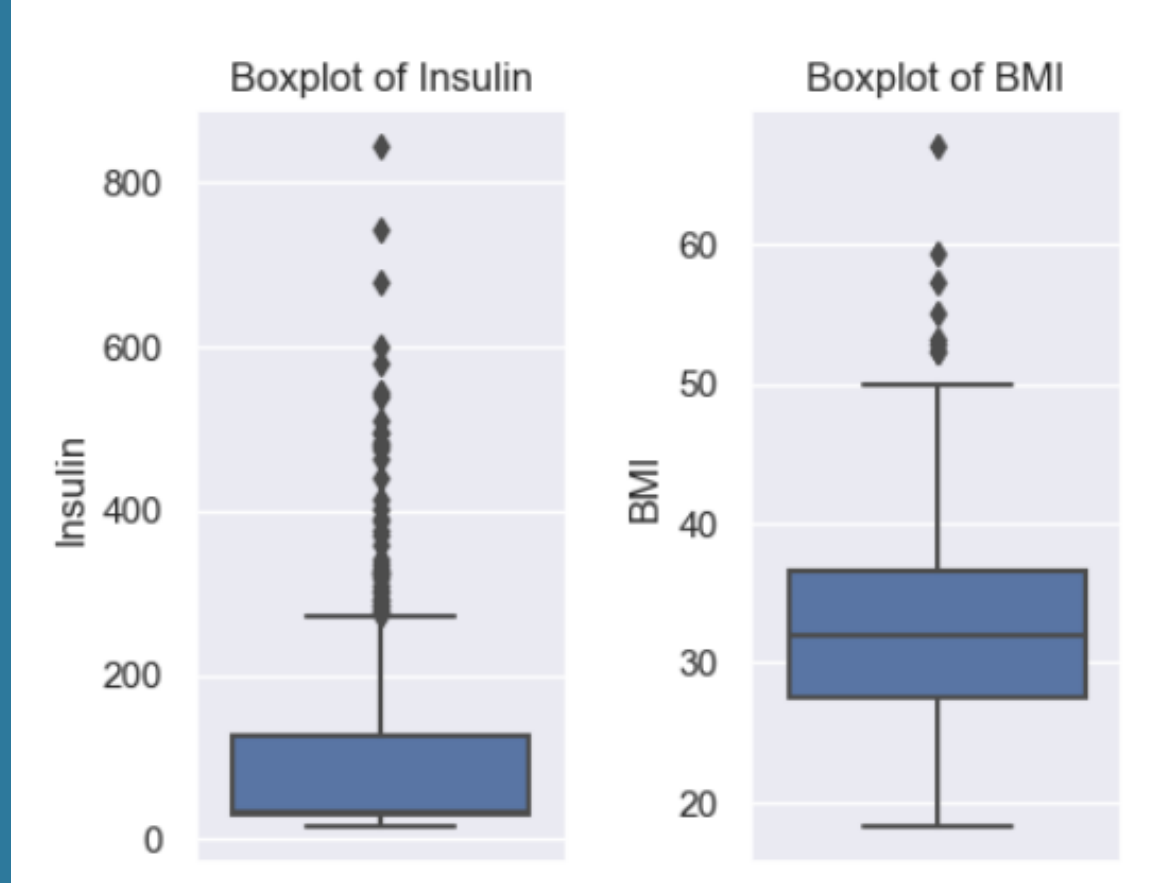
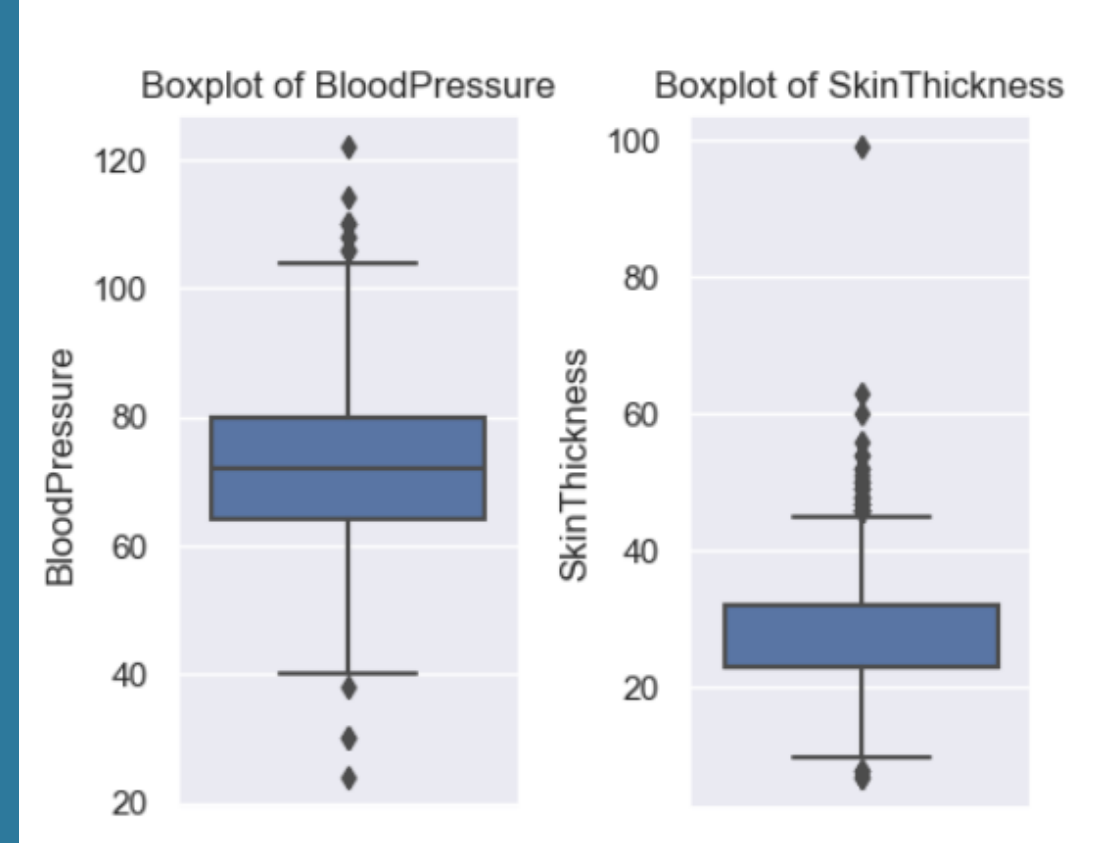
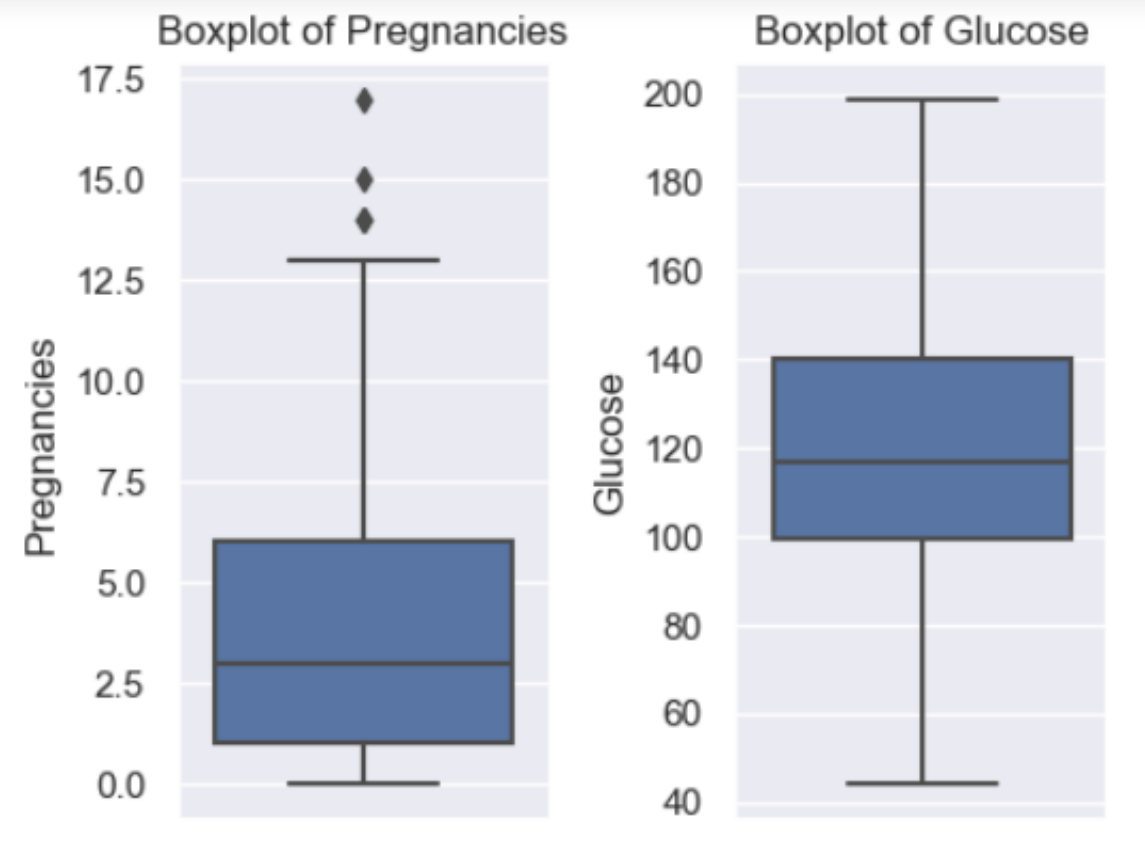
QQ
plots

Heatmap

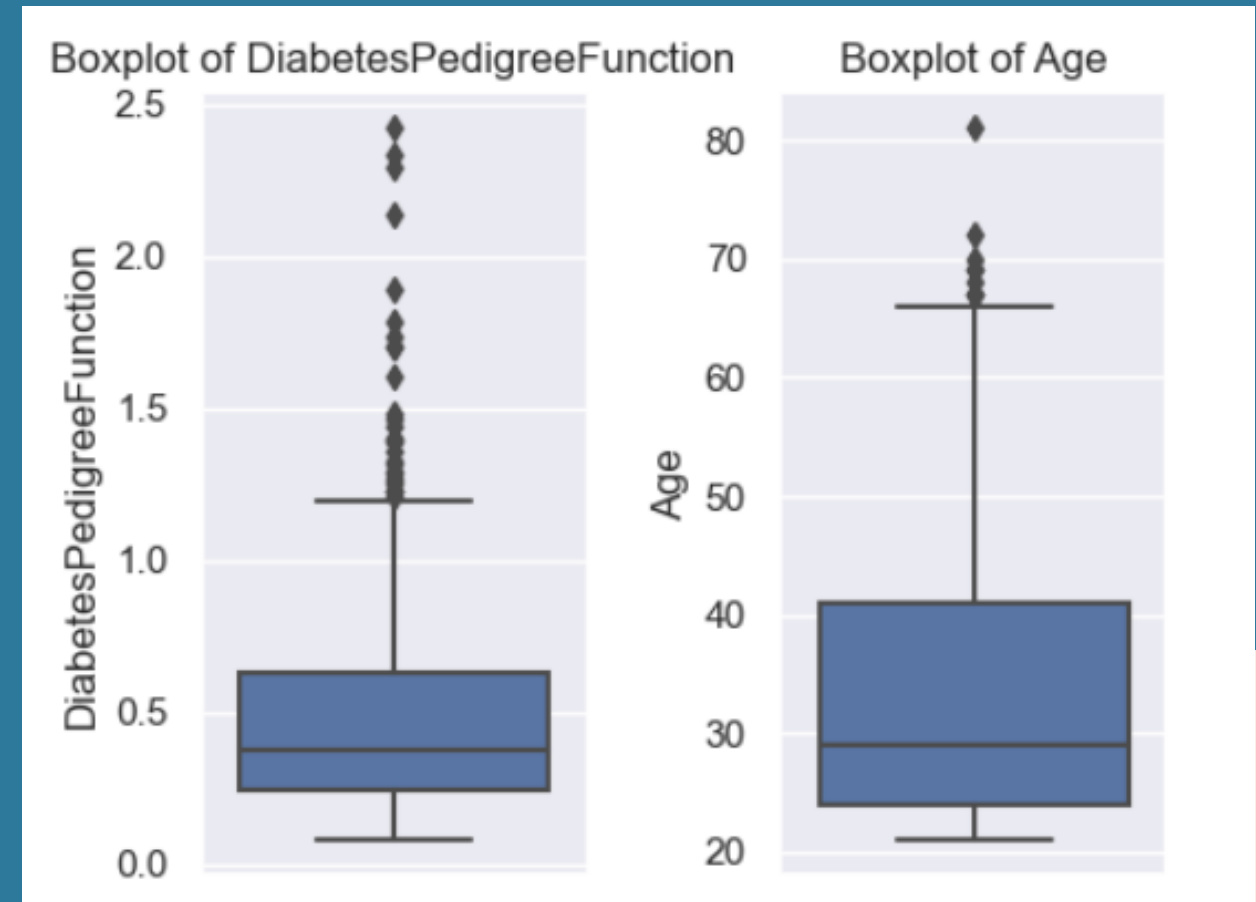


Checking the Balance of our data





Checking outliers:
Glucose is the only
feature without any
outliers



4. Building models:

1. Logistic regression

```
# One-hot encode categorical variables
diabetes_scaled = pd.get_dummies(diabetes_scaled, columns=['BMI_Category', 'Age_Group'], drop_first=True)

# Logistic Regression

X_lr = diabetes_scaled[['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI',
                        'DiabetesPedigreeFunction', 'Age', 'Glucose_Insulin_Interact',
                        'BMI_Category_Normal Weight', 'BMI_Category_Overweight', 'BMI_Category_Obese I',
                        'BMI_Category_Obese II', 'BMI_Category_Obese III',
                        'Age_Group_30-40', 'Age_Group_40-50', 'Age_Group_50-60', 'Age_Group_60-70', 'Age_Group_70+']]
y_lr = diabetes_scaled[['Outcome']]

X_train_lr, X_test_lr, y_train_lr, y_test_lr = train_test_split(X_lr, y_lr, test_size=0.2, random_state=42)

model_lr = LogisticRegression(penalty="l2", max_iter=100, solver="lbfgs", random_state=42)
model_lr.fit(X_train_lr, y_train_lr.values.ravel())
y_estimated_lr = model_lr.predict(X_test_lr)

print("Logistic Regression:")
print("Accuracy: ", accuracy_score(y_test_lr, y_estimated_lr))
print("\nConfusion Matrix: ", confusion_matrix(y_test_lr, y_estimated_lr))
print("\nClassification report: ", classification_report(y_test_lr, y_estimated_lr))
```



accuracy

score:

0.77

2. Random forest

```
n_estimators_rf = [20, 30, 40, 50, 70, 80, 90, 100, 120]
max_depth_rf = np.arange(10, 30)
accuracy_matrix_rf = np.zeros([len(n_estimators_rf), len(max_depth_rf)])

for i, estimator in enumerate(n_estimators_rf):
    for j, depth in enumerate(max_depth_rf):
        rf_model = RandomForestClassifier(n_estimators=estimator, max_depth=depth, random_state=42)
        rf_model.fit(X_train_lr, y_train_lr.values.ravel())
        y_estimated_rf = rf_model.predict(X_test_lr)
        accuracy_matrix_rf[i, j] = accuracy_score(y_test_lr, y_estimated_rf)

best_accuracy_rf = np.max(accuracy_matrix_rf)
best_params_rf = np.unravel_index(np.argmax(accuracy_matrix_rf), accuracy_matrix_rf.shape)
best_n_estimators_rf = n_estimators_rf[best_params_rf[0]]
best_max_depth_rf = max_depth_rf[best_params_rf[1]]
```




Best accuracy
score: 0.79



Best accuracy
score: 0.78



Heatmap of accuracy score for AdaBoost



3. AdaBoost

```
# AdaBoost Grid Search
```

```
base_classifier = DecisionTreeClassifier(max_depth=2)
n_estimators_ab = [20, 30, 40, 50, 70, 80, 90, 100, 120]
learning_rate_ab = [0.01, 0.02, 0.05, 0.1, 1]
accuracy_matrix_ab = np.zeros([len(n_estimators_ab), len(learning_rate_ab)])

for i, estimator in enumerate(n_estimators_ab):
    for j, rate in enumerate(learning_rate_ab):
        adaboost_class = AdaBoostClassifier(base_estimator=base_classifier,
                                             n_estimators=estimator,
                                             learning_rate=rate,
                                             random_state=42)

        adaboost_class.fit(X_train_lr, y_train_lr.values.ravel())
        y_estimated_ab = adaboost_class.predict(X_test_lr)
        accuracy_matrix_ab[i, j] = accuracy_score(y_test_lr, y_estimated_ab)
```


4. LightBGM

```
train_data=lgb.Dataset(X_train_lr,label=y_train_lr)
test_data=lgb.Dataset(X_test_lr,label=y_test_lr,reference=train_data)

params={"objective":"binary",
        "metric":"mse",
        "boosting_type":"gbdt",
        "num_boost_round":100,
        "learning_rate":0.05,
        "max_depth":6,
        "num_leaves":30,
        "feature_fraction":0.9}

# train model
lgb_model = lgb.train(params,train_data,valid_sets=[test_data])


# make predictions
predictions=lgb_model.predict(X_test_lr,num_iteration=lgb_model.best_iteration)
binary_predictions = [1 if pred>=0.5 else 0 for pred in predictions]

# evaluate the model
lbg_accuracy=accuracy_score(y_test_lr,binary_predictions)
print("Accuracy of LightBMG model is ",lbg_accuracy)
```

**accuracy
score: 0.76**

5. SVM

```
Best SVM model parameters:  
{'C': 0.1, 'kernel': 'linear'}  
The best SVM model accuracy is 0.81233333333333335
```



6. XG Boost Classifier

Best XGBoost model parameters:

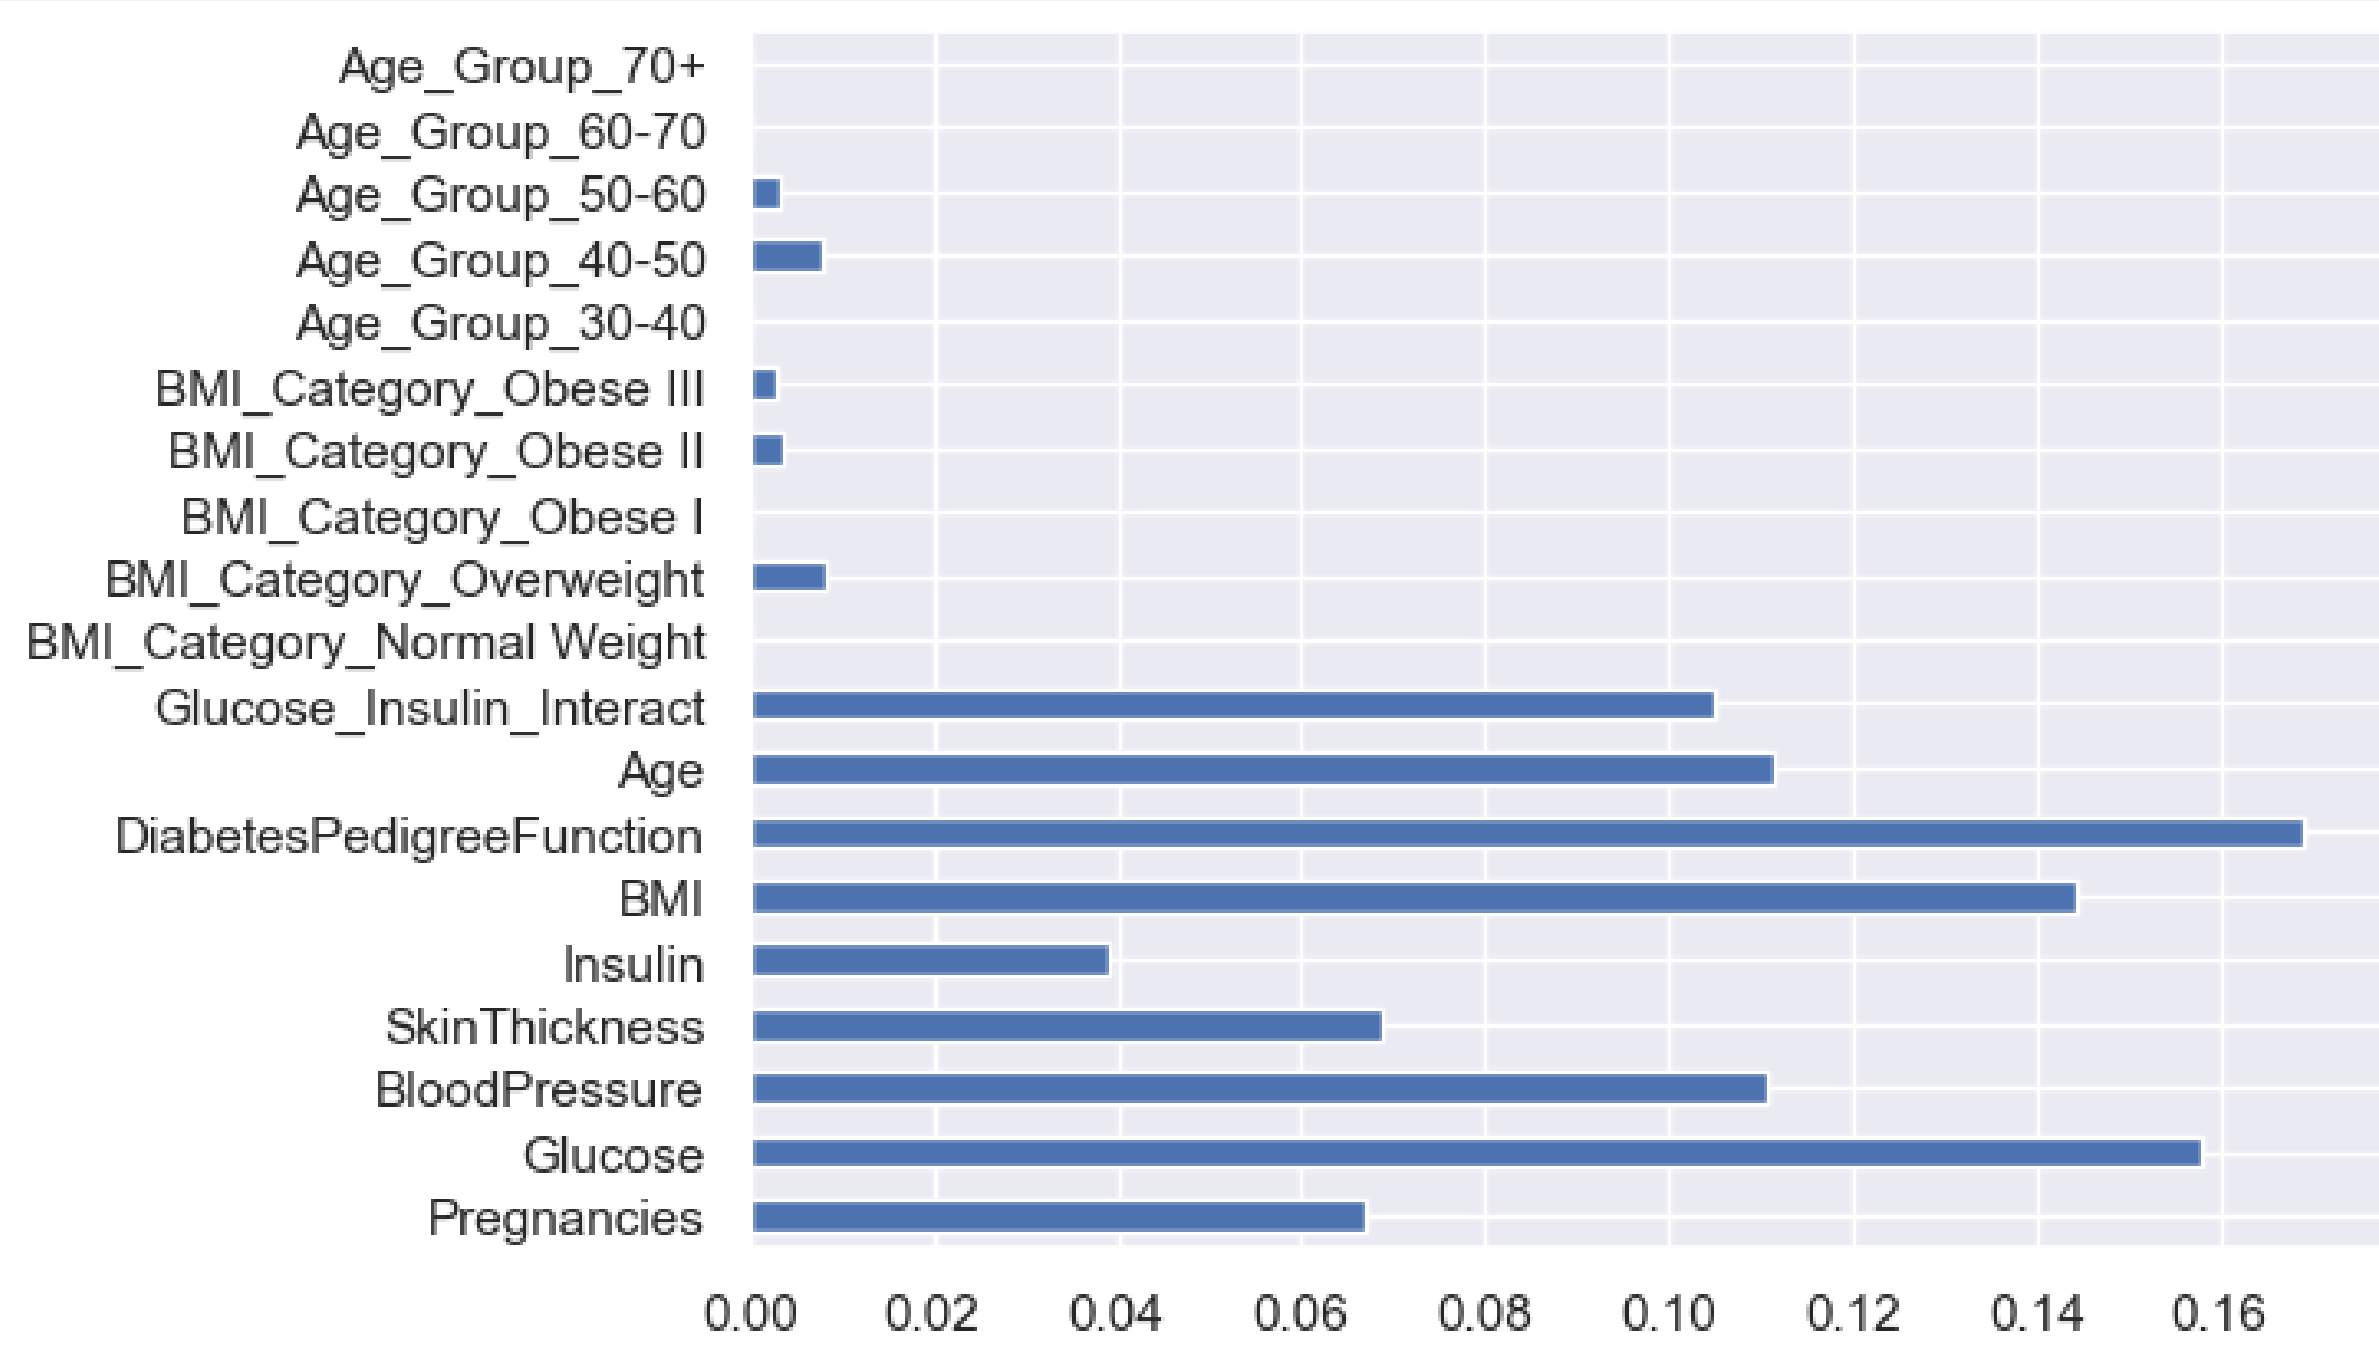
```
{'colsample_bytree': 1.0, 'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 50, 'subsample': 0.8}
```

The best XGBoost model accuracy is 0.783379981340797

**Accuracy
Matrix**

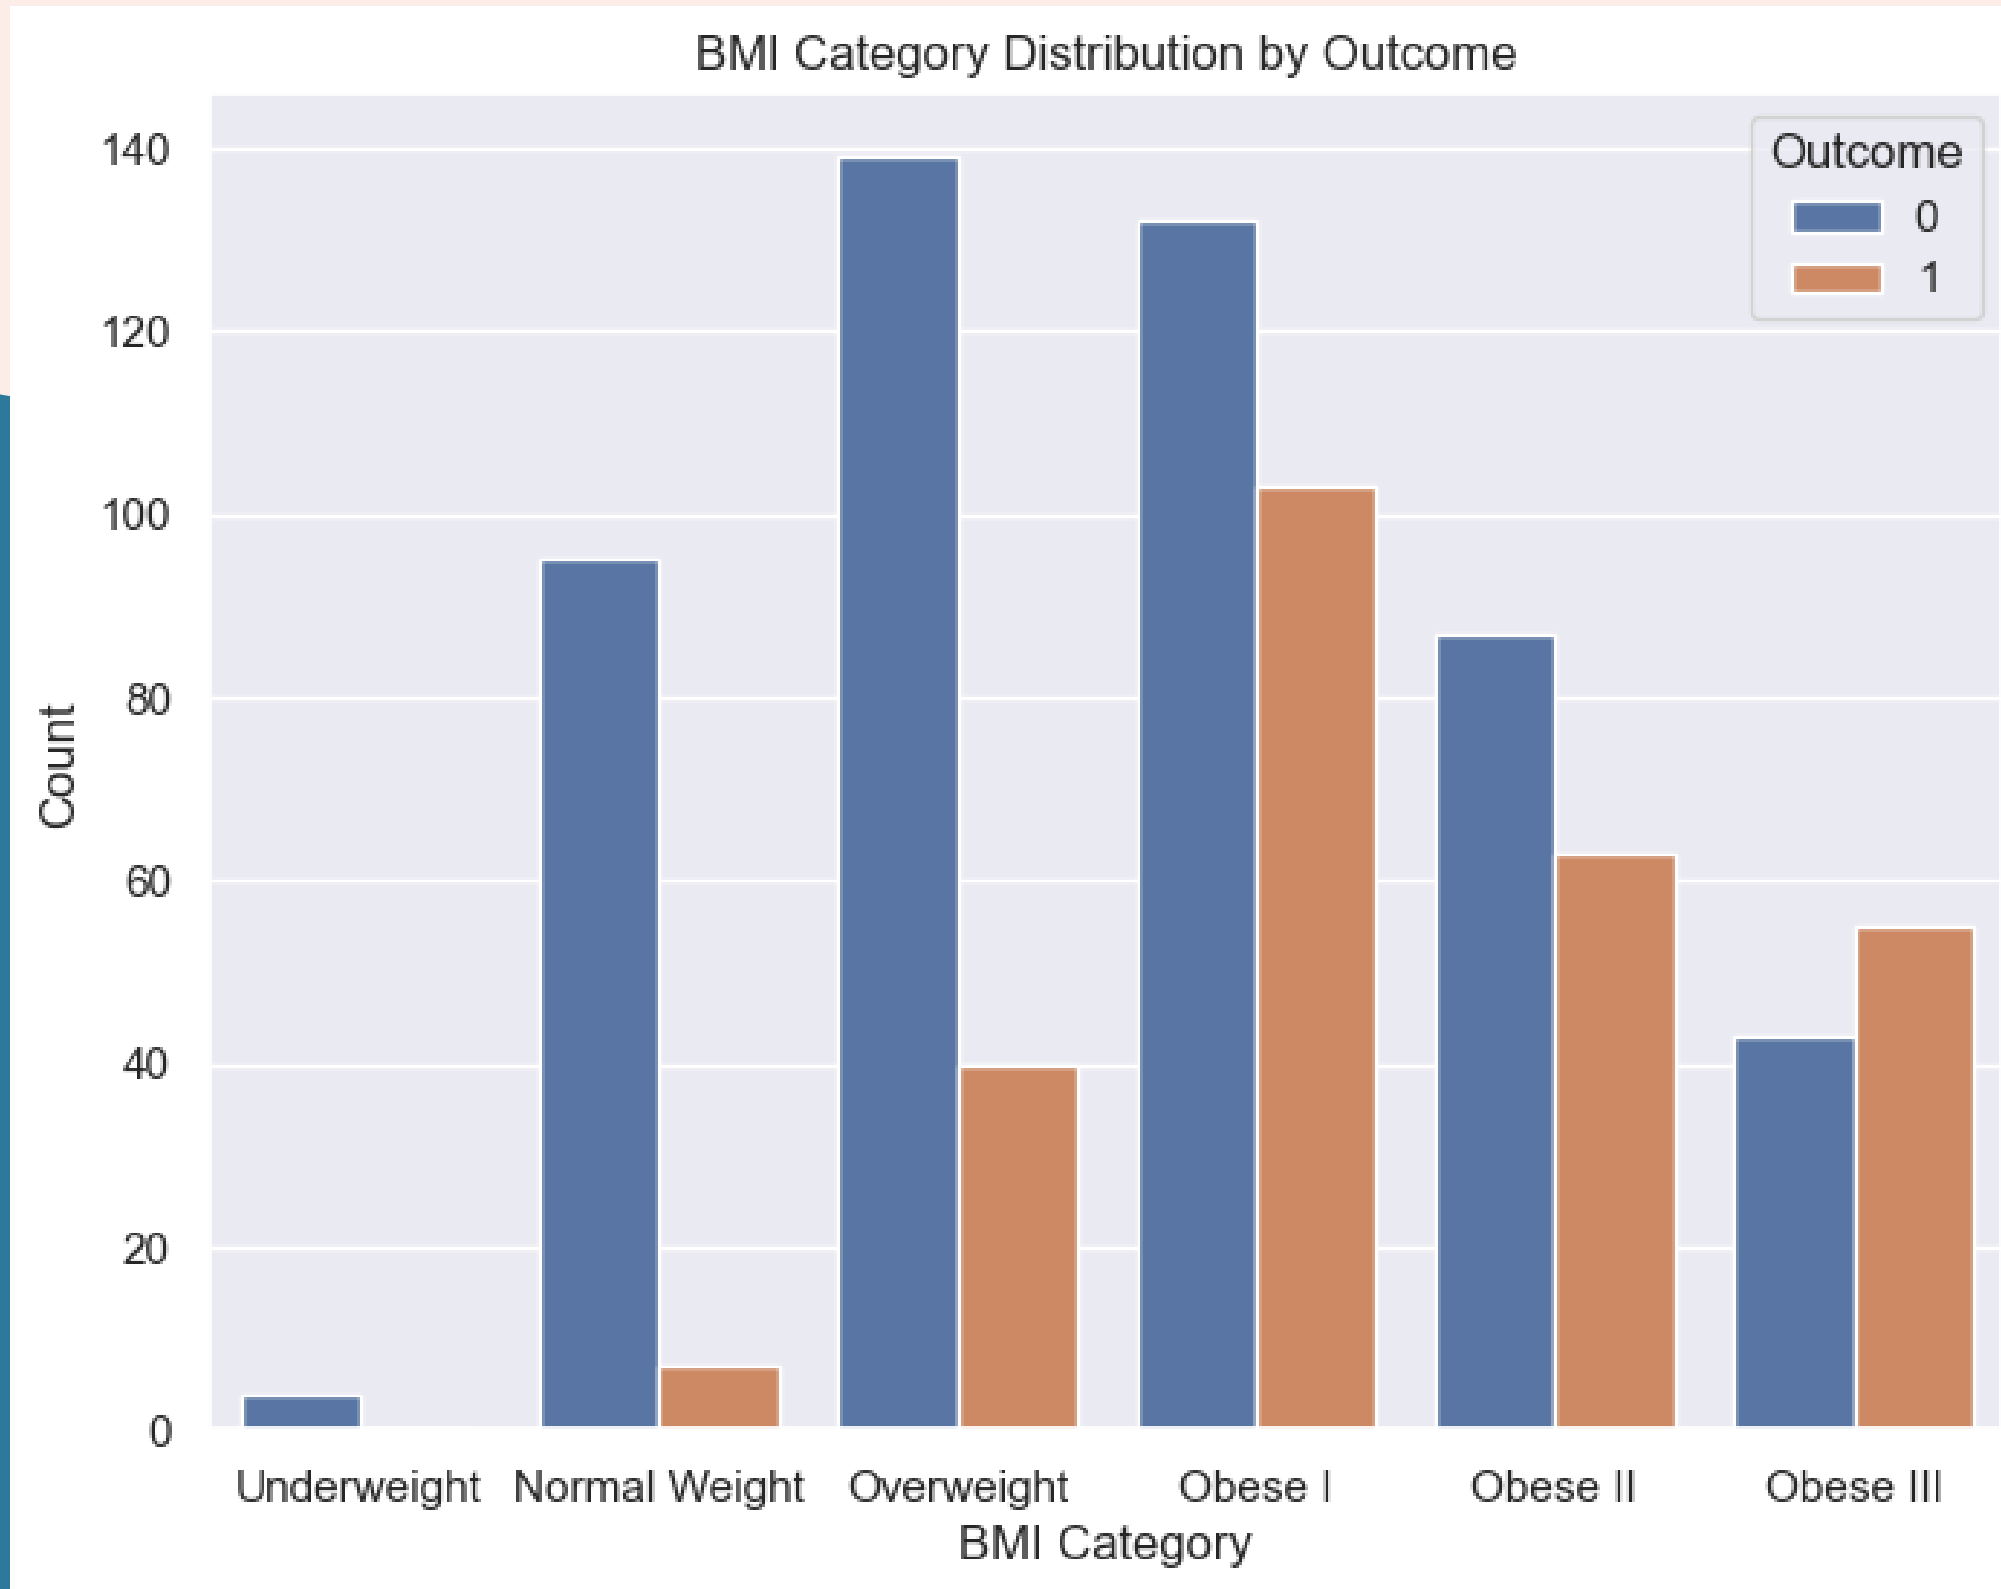
```
[[0.70030654 0.70683727 0.75737705 0.73779821 0.76871918 0.7622551
 0.70686392 0.7100893 0.76060243 0.74435559 0.76222844 0.7605891
 0.71338131 0.7198454 0.75409836 0.73775823 0.76544049 0.74919366
 0.78175397 0.78010129 0.76871918 0.76220179 0.75731041 0.7638278
 0.75731041 0.7703452 0.77361056 0.76378782 0.75893643 0.77031854
 0.75572438 0.75888311 0.75732374 0.76867919 0.75408503 0.76705318]
[0.75893643 0.76707983 0.76545382 0.76871918 0.76218846 0.75565774
 0.76222844 0.75729708 0.75896308 0.76381447 0.75241903 0.74260962
 0.75728375 0.76057577 0.75403172 0.76546715 0.75404505 0.75893643
 0.70848994 0.72311076 0.76222844 0.75087298 0.77359723 0.7654938
 0.7117553 0.71824603 0.75733707 0.74270292 0.76869252 0.75572438
 0.71339464 0.72152472 0.7605891 0.73451953 0.76869252 0.74594162]
[0.78337998 0.77527656 0.76548047 0.77198454 0.76221511 0.76380115
 0.76709316 0.76709316 0.76057577 0.77031854 0.76545382 0.76705318
 0.7621618 0.75728375 0.76381447 0.75732374 0.76545382 0.76709316
 0.7621618 0.76875916 0.76214847 0.7621618 0.76218846 0.76706651
 0.77029188 0.76871918 0.76381447 0.7638278 0.75563108 0.76221511
 0.7589231 0.7638278 0.75405838 0.76054911 0.7589231 0.75565774]]
```

Feature Importance



Features that make most contributions are DiabetesPedigreeFunction, Glucose, BMI, Age, BloodPressure.

Feature Engineering

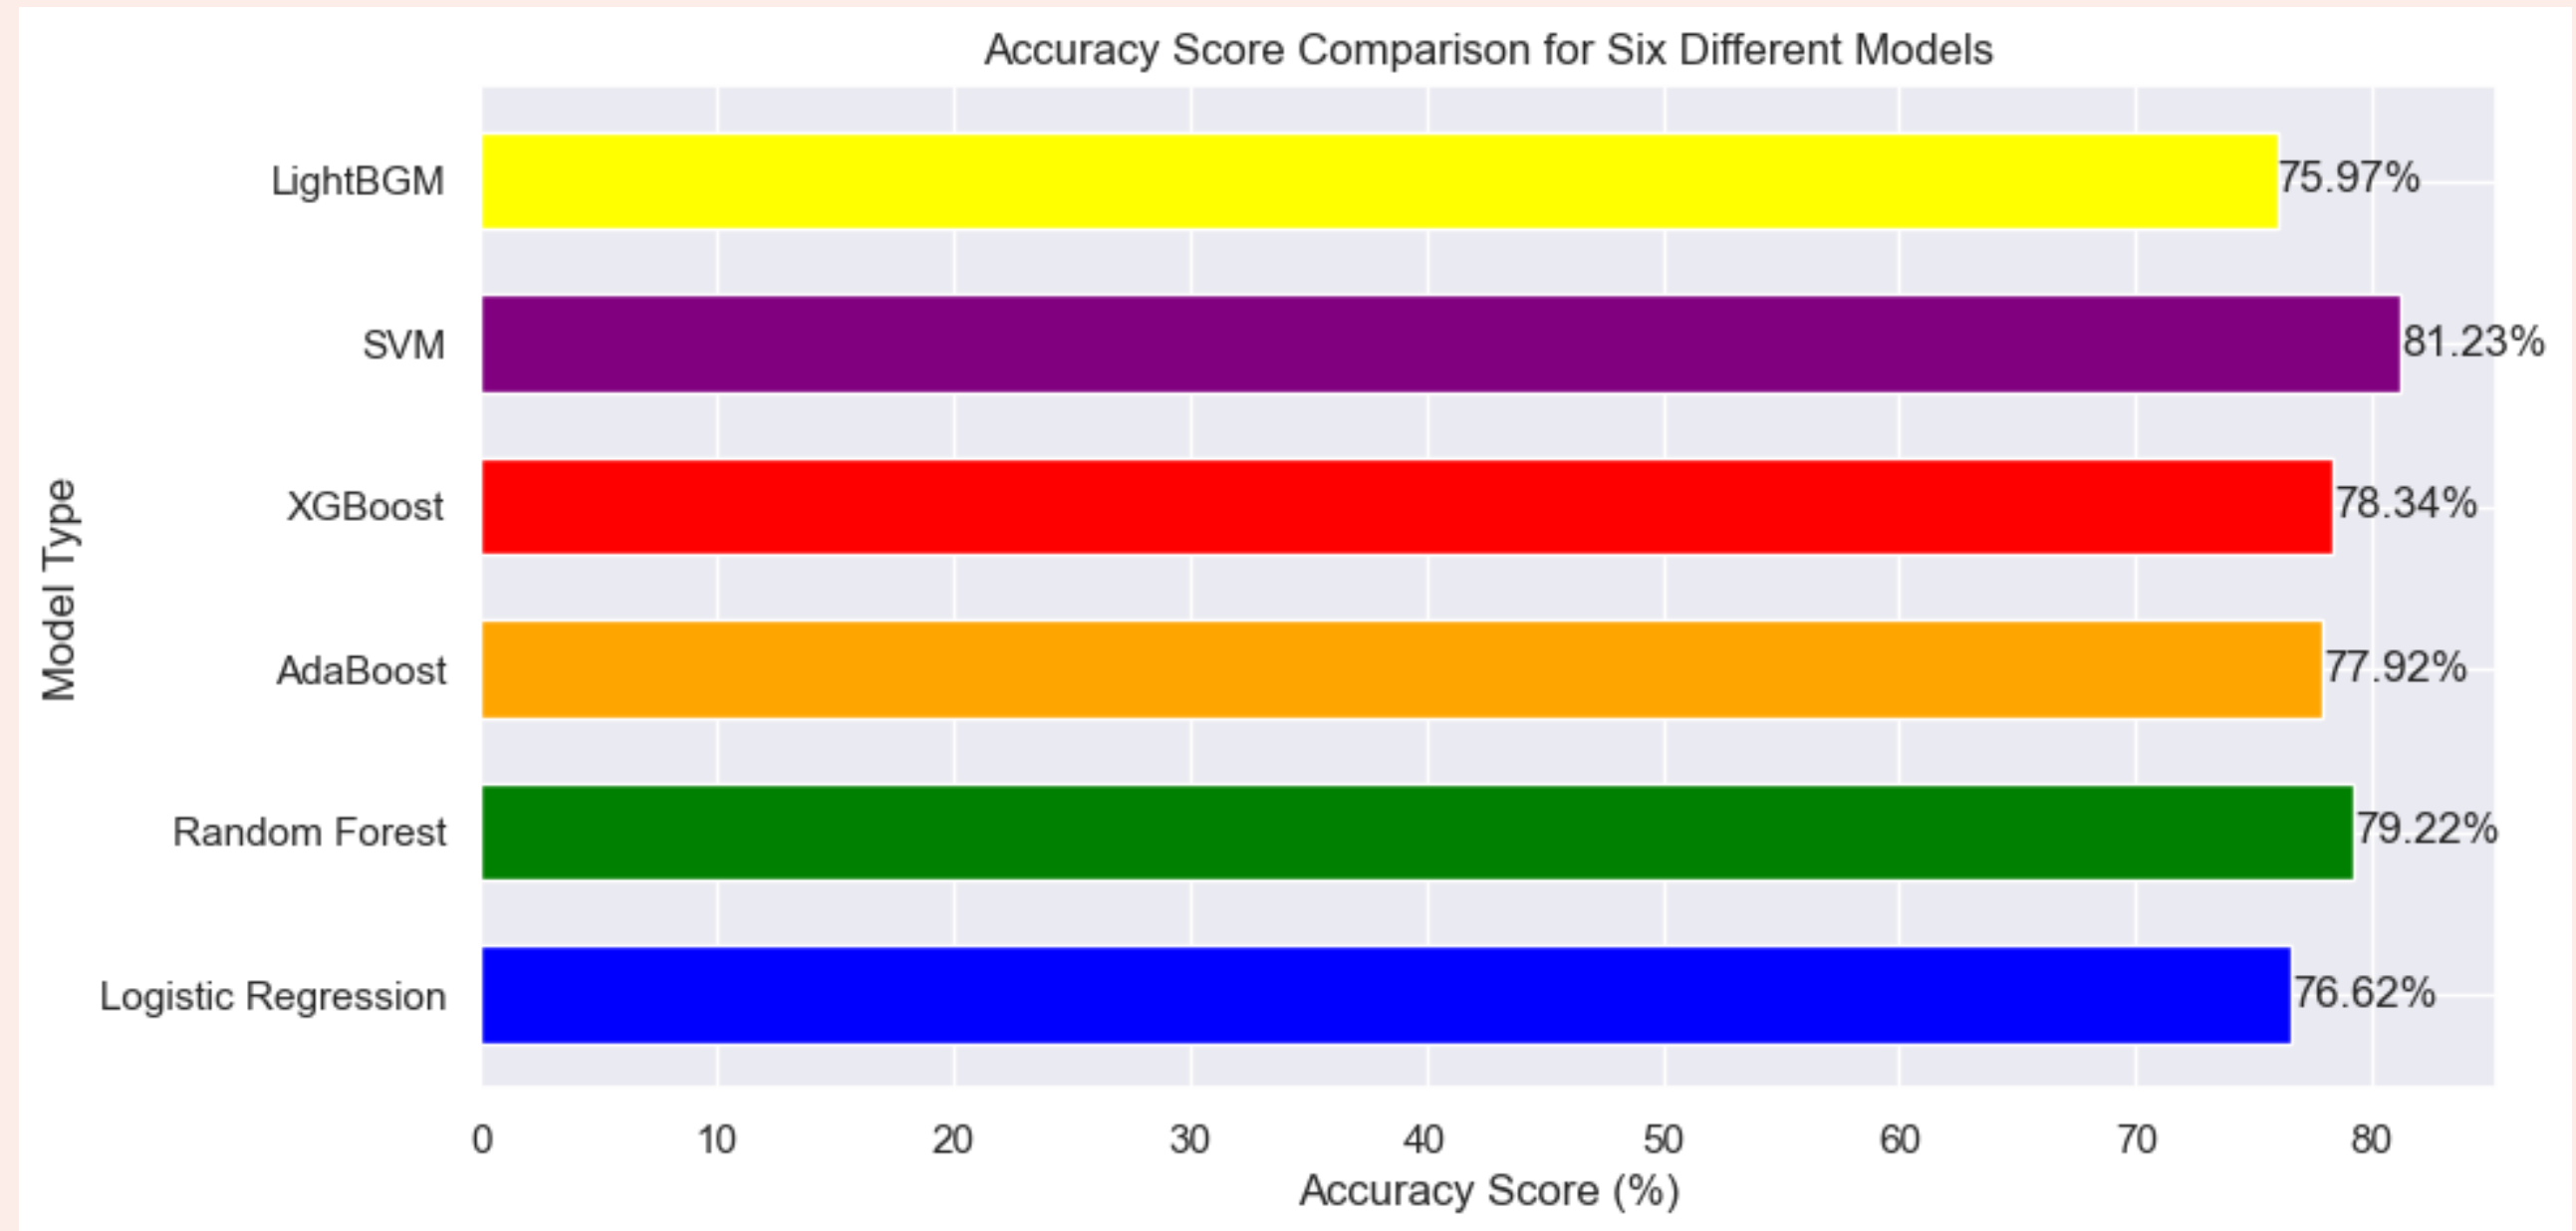


The following new features were added:

- BMI_Cateogy
- Glucose_Insulin_Int
eract
- Age_Group

Performance Comparison

And the winner is!!!



Further Improvements

- **Better Feature Engineering**
- **Regularization**
- **Better Scaling of Data**

**THANK YOU FOR YOUR
ATTENTION!
WE ARE NOW OPEN FOR
ANY QUESTIONS!**

